

Sean Kim

909-618-8586 | skim8705@gmail.com | linkedin.com/in/seankim08 | github.com/sean-kim05 | sean-kim05.github.io

EDUCATION

University of California, Irvine

B.S. in Computer Science

Irvine, CA

Expected Jun 2027

EXPERIENCE

Software Engineer Intern

Jun 2025 – Sep 2025

MIADVG

Los Angeles, CA

- Shipped and maintained 3 React/TypeScript client web apps with Tailwind CSS, delivering features to 10,000+ combined users in 2-week sprint cycles
- Built REST endpoints with Express on Node.js across PostgreSQL and MySQL, unifying data access with typed error handling and consistent schemas
- Automated CI/CD on AWS EC2 via GitHub Actions, running build, test, and deploy on every push and shipping 2-3 releases weekly

Software Developer Intern

Jan 2025 – May 2025

CALIT2, University of California, Irvine

Costa Mesa, CA

- Developed a Python Flask computer vision system streaming detections from 5+ GPU-accelerated cameras to a React security dashboard at 100-200ms latency
- Fine-tuned YOLO with PyTorch transfer learning on client footage, improving mAP by 8 points over the stock baseline for production imaging conditions
- Implemented a threat-flagging feature in the security dashboard that flagged high-severity events (fires, vehicles, suspicious activity), reducing redundant alerts for the security team

Undergraduate Research Assistant

Jan 2024 – Jun 2024

University of California, Irvine

Irvine, CA

- Built Python pipelines that extracted and cleaned thousands of UCI documents from PDF and HTML into training-ready text for iGPT (now ZotGPT)
- Tested and evaluated the fine-tuned GPT-3.5 model on UCI-specific queries for ZotGPT, now serving 36,000+ undergraduates

PROJECTS

CollabCode [\[GitHub\]](#) | *React/TypeScript, Monaco, Flask-SocketIO, Redis, PostgreSQL*

Feb 2026 – Apr 2026

- Built a real-time collaborative code editor (live cursors, follow mode, multi-file workspace) sustaining 50 concurrent clients/room at 17ms p95 sync latency
- Designed a three-layer persistence system (Redis cache, PostgreSQL source of truth, SQLite snapshots) backing live sessions and replay
- Integrated Claude API across four modes (generate, fix, explain, improve) with sandboxed server-side code execution

Briefly [\[GitHub\]](#) | *Cloudflare Workers, Durable Objects, Workers AI, Llama 3.3, JavaScript*

Apr 2026 – Present

- Engineered an AI standup tool on Cloudflare Workers with custom SSE streaming, forwarding Llama 3.3 tokens to the client live while capturing full responses server-side
- Implemented a Durable Objects backend with CRUD across personal and team-room namespaces, plus an endpoint generating weekly summaries via a chained second LLM call

DSA Visualizer [\[GitHub\]](#) | *React, Vite, Flask, Anthropic SDK, Python, CI/CD*

Jan 2026 – Feb 2026

- Developed an interactive visualizer covering 12 algorithms across pathfinding (Dijkstra, A*, BFS, DFS), sorting, BSTs, heaps, and graphs
- Added a Claude AI tutor that streams context-aware explanations via SSE from a Flask backend

TECHNICAL SKILLS

Languages: Python, JavaScript, TypeScript, Java, C/C++, SQL, HTML/CSS

Frameworks & Libraries: React, Node.js, Express, Flask, Tailwind CSS, PyTorch, NumPy, Pandas

Cloud & Infra: AWS (EC2), Cloudflare Workers, Docker, GitHub Actions, CI/CD, REST, Vite

Databases: PostgreSQL, MySQL, Redis (NoSQL), SQLite

AI/ML: Computer vision, LLM integration, Fine-tuning, Transfer learning, Anthropic SDK